

A Tool for Visualizing Patterns of Spreadsheet Function Combinations

Justin A. Middleton, Emerson Murphy-Hill
North Carolina State University
Raleigh, North Carolina, USA

Abstract—Spreadsheet environments often come equipped with an abundance of functions and operations to manipulate data, but it can be difficult to understand how programmers actually use these in practice. Furthermore, users can combine several functions into a complex formula, complicating matters for both researchers and practitioners who want to study formulae to improve spreadsheet practices. Therefore, we developed a tool that visualizes patterns of function combination in spreadsheets as an interactive tree of possibilities. Using spreadsheets from the public spreadsheet corpora, we then apply to the tool to real datasets and demonstrate its ability to capture the most common and most anomalous patterns of function combination and their contexts in actual workbooks.

I. INTRODUCTION

Spreadsheets surround us, ordering the ever-growing flood of data that we produce. They serve key roles in industry, where as much as 95% of U.S. financial firms use them daily [1], and academia, where teachers can use them to track research and students' grades, not to mention many other lines of work [2]. It's this flexibility that grants them their allure: while a novice can still work without much experience, spreadsheets offer hundreds of specialized functions to fulfill a wide range of advanced needs, from complex statistics to text manipulation [3]. As such, it should be no surprise that tens of millions of workers rely on these programs every day [4].

However, this ubiquity is not without danger: as more people adopt these tools and make larger and more complex sheets, the cost of failure grows too. To stress this, the European Spreadsheet Risks Interest Group (EuSpRIG) documents many horror stories of spreadsheets gone terribly wrong¹. One such story tells of a 2013 economics paper which drew stark connections between debt and national growth. Its findings shaped many political initiatives throughout the United States and Europe. However, other researchers soon discovered that a critical coding error in the research spreadsheets hampered the results and that, by correcting this error, they completely reversed the findings! And this is only one of several accounts; even if most spreadsheet errors are relatively benign, poor spreadsheet practice can risk millions, if not billions, of dollars [5].

From these stories, we can see the dire importance of understanding how people actually use spreadsheets, if only to avoid disasters. Fortunately, spreadsheet researchers have

supported this by amassing spreadsheet collections from diverse sources, industrial to academic. Collections like EUSES [6], Fuse [7], and the Enron corpus [8] have already fueled productive research across the field, such as in Hermans' [9] or Jansen's [10] search for code smells² in spreadsheet environments or Aivaloglou and colleague's work on making a grammar to capture every possible spreadsheet program in Excel [11]. Additionally, some research has focused simply on comparing these collections, as Jansen did in 2015, to discover variation in spreadsheet techniques over time and location [12].

Considering the sheer value of spreadsheets and all of these resources to describe their use, it becomes necessary that we design tools to make sense of it all. Tools, then, that empower people to explore datasets are rife with the potential to improve the current design of spreadsheets themselves and educate new generations of spreadsheets users on how to work best.

This paper contributes such a tool, Perquimans?!, that attempts to fulfill this by building on the aforementioned spreadsheet corpora. Working with spreadsheets in Excel, Perquimans?! focuses primarily on the functions, like SUM and IF and hundreds more, which are the building blocks of formula and programs which determine values across a spreadsheet. The tool creates an interactive, exploratory visualization which captures how people combine these functions in a selected collection of spreadsheets, quantifying the broad patterns of function nesting while also being specific enough to find anomalous formulae and link them back to the specific cell in spreadsheets where they originate.

Alongside the tool, we discuss several of the key decisions we faced during design, such as what to do with redundant formulae in a sheet or how to represent optional arguments, and we describe how our overall design goals both helped and hindered this process. We then apply the tool to a number of case studies to evaluate how much it could contribute to several different contexts:

- API improvement
- Detection of code smells
- Spreadsheet education

To round out that discussion, we also conducted a user study to judge how usable and intuitive the tool actually is. This grounds the discussion of its potential uses while also

²Or, instances of inelegant code that do not necessarily cause problems alone but nevertheless suggest a need for refactoring to prevent future problems. More on this in section IV.

¹<http://www.eusprig.org/horror-stories.htm>

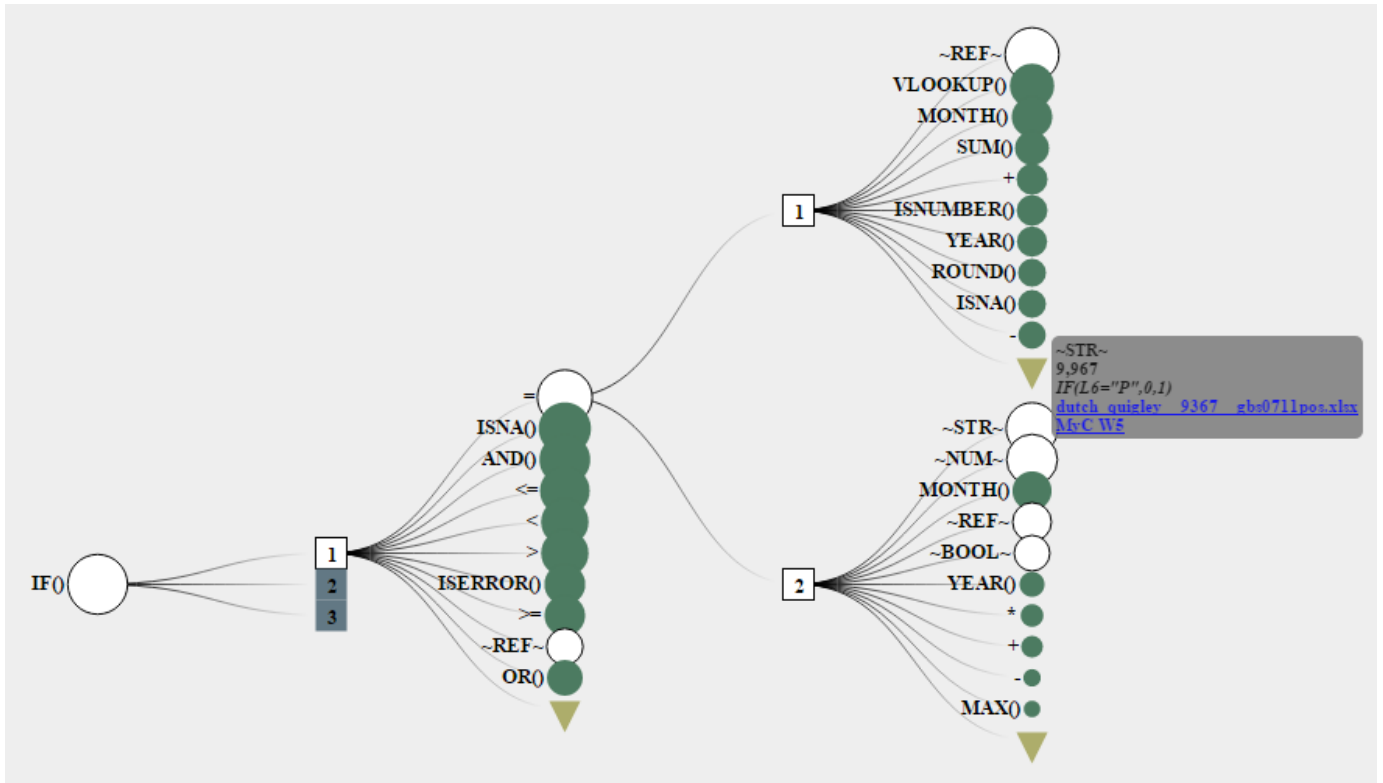


Fig. 1. A picture of Perquimans?!

underscoring some of its inherent limitations and setting the way for future extensions and improvements upon the concept.

II. RELATED WORK

Given humanity's affinity for visual information, visualizations have long been prized for their ability to communicate information quickly and efficient [13]; the visualization of spreadsheets, therefore, has offered much room for research. Several previous tools work within a spreadsheet, seeking to clarify the connections between the visible values in cells, the formulae that produce them, and the dependencies between cells that inform them. Igarashi and colleague's fluid visualizations [14], for example, do this by imposing lines, color, and animation over the spreadsheet to visually explain these tangled connections. Likewise, Clermont explored ways of grouping cells by color and border that were related through similar functions, neighbors, or references [15], and Hipfl extended this work by incorporating the layouts and labels of spreadsheets [16]. Other tools focus on creating visualizations external to the sheets. Hermans and colleagues address a spreadsheet programmer's information by making dataflow diagrams from individual sheets through the tools GyroSAT [17] and Breviz [18]. Perquimans?! is among the latter class, creating standalone visualizations about spreadsheets. In contrast to Hermans' work, though, our work focuses on representing functions and how they comprise complex formula, rather than exploring the data structures of any particular spreadsheet.

Many other tools inspect the structure of formulae but often from the perspective of finding bad design. Following Fowler's initial description of code smells in object-oriented design [19], many researchers have applied similar observations to spreadsheets and have cataloged classes of problematic formulae [9] [20] [21], often accompanying them with detection tools [22] and investigations [10]. For example, Hermans, leveraging her work with the aforementioned dataflow diagrams, has worked with assessing common inter-worksheet smells, such as feature envy and inappropriate intimacy, wherein references between worksheets (rather than cells) in a file suggest a problem. [23]. Other tools, like Badam's Excel refactoring [24], are concerned more with understanding the formula in the service of changing them. However, these studies define standards of poor design and orient their searches around them. Though Perquimans?! supports the search for bad design, as we will see, the tool itself takes a much more agnostic approach in presentation, offering examples without regard for design quality.

Implicit in this study is an exploration of the Excel API and how people use it. Previous research, such as the first study on the Enron spreadsheets [8] or in its comparison with the EUSUS corpus [12], quantify how often users employ certain functions but not how they combine them. Researchers have applied such questions to other contexts as well, such as Murphy and colleague's study into how Java developers use Eclipse and what commands they use most often [25]. Others focus on the documentation of the API itself: Robillard

and DeLine, for example conducted a series of surveys and interviews to diagnose common problems with learning APIs, and they found that code examples, one of the focuses of this tool, is one of the most important aids to have in learning a new system [26]. Though this paper applies the tool to the Enron corpus specifically, the tool itself is not limited to this; it can be readily applied to any collection of Excel spreadsheets a user gives it.

III. APPROACH

A. Design Goals

In visualizing the spreadsheet data, we outlined the core goals for what the tool should accomplish. Furthermore, with these goals, we also considered the likely challenges of achieving those goals and particular outcomes which we considered beyond the scope of the research. We present both sides of these goals below:

- 1 **Provide an interactive environment for free exploration.** There are a lot of data in some of these collections, and the user should be able to pick for themselves which to explore while ignoring the rest. Furthermore, this exploration should not be designed around answering one specific question but rather be open-ended.
- 11 **Do not make judgments about good and bad design.** Though users may certainly approach the tool with the intention of finding particular formulae, Perquimans?! will not prioritize any on metrics of good or bad practice. Rather, it will present them agnostically, referring only to frequency.
- 2 **Emphasize the quantitative patterns in formula construction.** Perquimans?! should address the questions of how often the spreadsheet programmers used a certain function and where they used it. In this way, it must show precise metrics, such as frequency of use and depth of function nesting, of the dataset it conveys.
- 12 **Do not overwhelm with numbers.** For popular functions especially, we anticipated some particularly dense trees. Though everything should be accessible in accordance with Goal 1, we should design in a way that reduces noise.
- 3 **Promote a qualitative understanding of the patterns.** Perquimans?! should supplement its observations with concrete instances of relevant formulae from the corpus. Ideally, it should even direct users to the exact cell in the spreadsheets where the formula was used, contextualizing the functions.
- 13 **Do not attempt to explain formulae.** Though Perquimans?! tries to foster understanding by linking pattern to example, it should not provide a precise description of what a complex formula accomplishes. The tool's user must infer this on their own.

Though we kept these guidelines in mind through the development of Perquimans?!, we nevertheless reached a number of points in which these principles conflicted. For a discussion of these decisions, see section III-E.

B. Walkthrough of the Tool

Before discussing the minutiae of tool's development, though, we will give a basic example of how to use the tool. For this, we will answer this question: *What kinds of functions do people use to define the condition in an IF function?* The full visualization that answers this is available in figure 1; however, to understand the interaction, we will start from the beginning.

First, a word on the IF function. In Excel, it serves as one of the logical functions: given a condition as the first argument that evaluates to either true or false, it takes on the value of either the second or third argument, respectively. Or, as the documentation writes out, "IF(Something is True, then do something, otherwise do something else)"³. Furthermore, the third argument is optional; if left out, the function will return a default "FALSE" value if necessary. Examples of this function in practice can be seen in III-B.

| | A | B | C |
|---|----|--------------------------------|---------------|
| 1 | | Formula | Result |
| 2 | 10 | =IF(A2 >= 10, "True", "False") | True |
| 3 | | =IF(A2 < 10, "True", "False") | False |
| 4 | | =IF(A2 >= 10, "True") | True |
| 5 | | =IF(A2 < 10, "True") | FALSE |

Fig. 2. Configurations of the IF function

When the user first approaches the tree for a given top-level function – that is, a function nested within none other – only a few nodes are visible, as shown in figure III-B.

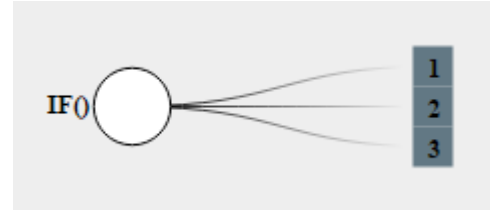


Fig. 3. How the IF tree looks at the start

The visualization so far comprises two types of node: the circle ●, which represents a discrete function in the formula and is size according to its frequency relative to the number of times to root function appeared⁴; and the numbered squares 1, which represent the positions of arguments within its parent function. From this, we can confirm that, of the times it was observed, IF can have at most three arguments passed into it, which corresponds with its specifications in the API⁵.

Knowing that it is the first argument which contains the conditionals, we click on the square labeled "1" to explore.

³<https://support.office.com/en-us/article/IF-function-69aed7c9-4e8a-4755-a9bc-aa8bbff73be2?ui=en-US&rs=en-US&ad=US>

⁴The root node, by definition, will always be the largest node in a tree.

⁵<https://support.office.com/en-us/article/IF-function-69aed7c9-4e8a-4755-a9bc-aa8bbff73be2>

To save space, when more than 10 unique arguments have been observed in any position, the tool displays only the first ten, with an option to display the rest by clicking the arrow

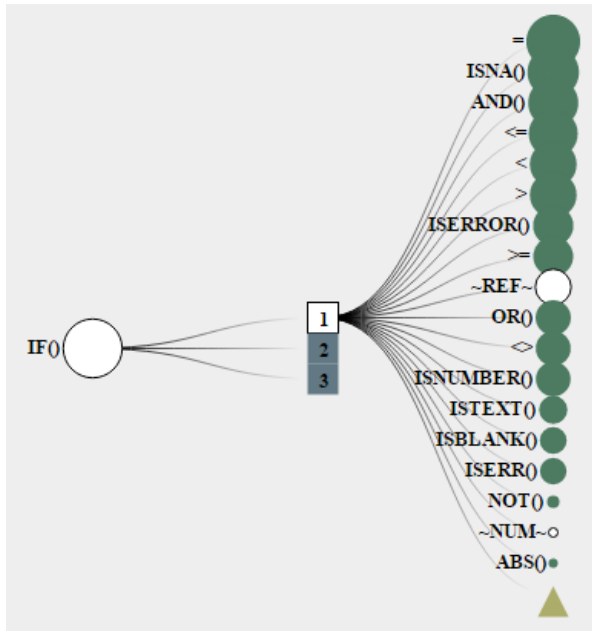


Fig. 4. First argument of the IF tree

We see that IF contains a number of comparison operators as its first argument, such as = and ≤, and boolean-returning functions, like ISNA and AND, with simple equality being the most common and some use as ABS being the least seen of everything actually used.

From here, we can further explore the common options among these functions. Clicking on the “=” node will yield two arguments and expanding each of them will peer into the range of common values of equality comparison, which Figure 1 shows.

For both sides of the equal sign, the operator has certain types of arguments which predominate over the others: on the left side is most often a reference to another cell; on the right, a string or number literal, which makes sense for the case of confirming a value in another cell before assigning this one. Furthermore, if the concept is difficult to imagine in practice, a tooltip accompanies each function node in the tree, providing a concrete example of a function that uses this structure and where it can be found. If this single instance is not enough, then the user also has the option to double-click the individual function node, which will open up a new tab with a table of many more examples from dataset.

C. Icon Glossary

From these images, we can provide a quick glossary of images within the visualization to aid interpretation.

- **ABS0** - **Function nodes** are green circles labeled with the function name. Their size is determined by their relative frequency, scaled logarithmically, to the frequency

of the root node in the tree. They may have any number of arguments, including zero, as determined by the API.

- **~BOOL~** - **Value nodes** are white circles with a generic label. They represent elements of a formula which accept no arguments, like references (A1), ranges (A1:A10)⁶, numbers, boolean values, strings, and errors. Functions with zero arguments, like NOW, however, should still be considered function nodes.
- **1** - **Argument nodes** are dark blue squares labeled with a number. The number is the position of an argument within the function in the parent node. In III-B for example, argument node 1 shows every function or value that was seen in the first argument.
- **Solid lines** represent a function’s ownership of its various arguments. As such, they connect the function nodes to argument nodes.
- **Fading lines** represent the possibility of an argument position to have several different possible functions or values. As such, they connect argument nodes to function or value nodes.
- **Expansion arrows** appear within an argument node when there are over 10 possible functions or values in an argument position. By default, only the first 10 will be shown along with this arrow; clicking this will show the rest as well.
- **1** - **Optional argument trees** are variations of trees that show only formulae with a specific number of arguments, allowing the user to control for the presence of optional arguments. By default, trees with dark blue nodes contain every type of formula without regard how many arguments they have – for example, single-argument ANDs could appear alongside 4-argument ANDs. See Figure III-C to see it visualized.




Fig. 5. The contents of each type of tree.

⁶That is, only if both sides of a references are cell references. If either side is a function that returns a cell reference, like INDIRECT, then the colon should be a function node.

D. Implementation Details

The final visualization is the product of two discrete processes:

- *Collection*: Given a set of Excel sheets, the tool, written in Java, uses Apache POI⁷ to identify and iterate over every cell containing a valid formula. Afterward, it calls POI's formula parser to break the formula text into an ordered set of individual tokens, which the tool then parses into the tree-like form. When all formulae have been analyzed like this, it produces JSON files for each top-level function in the set.
- *Presentation*: The JSON files, meanwhile, feed into the presentation code, implemented in Javascript with much help from the visualization library D3⁸. We chose D3 primarily for its accessibility: given the JSON, the visualization can be simply embedded into a webpage accessible through a browser.

A demonstration of the tool can be found at <https://github.com/DeveloperLiberationFront/Excel-Function-Visualizer>.

E. Design Decisions

Early in the design, we chose the tree form for its inherent ability to represent the loose parent-children relationship of formulae as functions and their arguments, which could be yet more functions. By adapting this structure to accommodate a broad range of possibilities for nested functions, the branching factor depends alternatively on the number of arguments in a function and the number of possible functions observed as an argument in a function.

Guided by the goals in section A, we faced a number of decision points, a sample of which we describe below, before we arrived at the design shown above.

- *Copied formulae*: Excel allows users to spread a formula over an area, repeating the same task in each cell with minor adjustments. Without checking for this, the analysis may not reveal the functions most commonly used together but rather the formulae most often applied to large areas. To combat this, we converted formulae from their native A1 format to the relative R1C1, in which copied formulae should be identical, and considered only each R1C1 formulae once per worksheet.
- *Importance of depth*: When a function appears within another, should it be analyzed only as a nested function or would it also be valid to analyze the nested function on its own?

For example, in the pictured IF function, we found that people often use another IF statement as the second or third arguments of the top-level IF. If we only consider functions exactly where they're found embedded in the formula, then information about the same function – IF, in this case – will be scattered across different trees with no way to aggregate them. If we record every instance of

a function by ignoring their context – that is, including an IF embedded within a SUM function in the same node as the top-level IF – then the tool will represent some functions in multiple nodes to capture every possible level of nesting. For now, we only include the first, with the latter to be added as an option in the future.

IV. CASE STUDY

The obvious question to pose to any visualization is this: Why does it matter that we see this way? In other words, it is not enough to prove that something can be visualized; instead we must also demonstrate what we can be gained from any particular visualization. As such, we raised a few tasks in the introduction in which we thought the tool could help: identifying places for API improvement, detecting bad smells or common programmer misunderstanding, and guiding spreadsheet education. To demonstrate the tool's applicability, then, we sought out places in the tree which best exemplify these respective concerns.

A. Bad Smells

Fowler's description of code smells [19] underscored an important point of code quality: between perfect code and bug-crippled spaghetti, there is a spectrum of code designs which, by themselves, are not faulty but nevertheless suggest problems and vulnerabilities in code design. Code smells, as we call them, soon commanded attention and earned wide-ranging applicability, eventually making it to the art of spreadsheet design.

Since then, spreadsheet researchers have taken the concept, if not the specific smells themselves, and applied them to spreadsheet programming [9] [20] [10] [21]. While some of the smells lie beyond the scope of the tool, such as those characterizing inter-worksheet connections, some have structures which create distinctive features within the tree, leading to easy detection through visualization:

- *Long Method*: Also called "Multiple Operations" in [9], this refers to situations where numerous functions and operators are used all within the same cell, creating a formula that tends to extend beyond a compact viewing space and makes it difficult to understand. In the visualization, since each nested and combined function creates yet another level of depth within the tree, these constructions can result in long, horizontal chains of nodes, as seen in figure IV-A.
- *Long Parameter List*: [tall, vertical towers of blue nodes]
- *Conditional Complexity*: [IF nodes within IF nodes]

Some examples are meaningful not for the space they take up but for how they clash with the expectations around a function. For example, Excel offers a number of functions that accept any number arguments given to them, up to a resource-defined limit. SUM and CONCATENATE are two such functions: adding more arguments simply adds or appends that element in the same way as every previous argument. Strange cases were found, then, in the cases where these functions were used with only one argument.

⁷<https://poi.apache.org/>

⁸<https://d3js.org/>




Fig. 6. The horizontal length of a tree may indicate complex or redundant design.

Otherwise, the LOOKUP functions are particularly notorious for their confusing mechanisms and, as such, have warranted some especial attention with regards to spreadsheet design. In a recent study, Hermans, Aivaloglou, and Jansen probed the Enron dataset for uses of VLOOKUP and HLOOKUP specifically. The issue at hand was the optional fourth argument: when omitted or set to TRUE, the function attempts to return a value from another column by matching two values approximately, but when set to FALSE, the match must be exact. Furthermore, exact matching works only with sorted columns or rows; if the programmer neglects to sort values before calling the *LOOKUP function, they may acquire inaccurate results [?].

Given access to the same dataset, how might this tool approach the issue? On the surface, the tool supports easy quantification through the tooltips displayed for each node, as shown in figure IV-A, answering the question, "How often do spreadsheet programmers use this function?". But, more to the point, the tool's distinction between the numbers of arguments for a single function lets the user compare frequencies of three-argument forms with the four-argument. The visualization does not, however, discern between the different values for booleans in the tree, much less whether the spreadsheet columns are ordered – but the linking of examples to patterns allows quick insight into which files are at risk for the problems they outlined.

Also at hand are issues with the standard LOOKUP function. Unlike the two variations above, LOOKUP accepts either two or three arguments, depending on which API-defined form the programmer uses: the vector form or the array form, the primary difference being whether the second parameter is a one-column/-row range or a two-dimensional array. The array form, as it happens, is strongly discouraged by the API in favor of the V/HLOOKUP, a notion that might also be leveraged in the pursuit in good spreadsheet design.

B. Education

Learning by example has long been a cornerstone of effective teacher techniques: by showing the student a concrete example of something, rather than dwelling in the realm of abstract description, comprehension increases [I'm sure there's a cite out there somewhere.] This is no different in software,

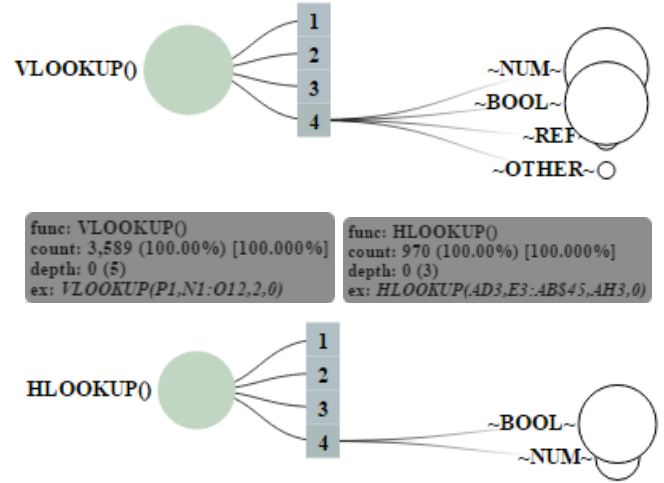


Fig. 7. VLOOKUP and HLOOKUP with four arguments. Numbers not expected to correspond perfectly with [?].

where every explanation of a tricky programming concept is much untangled by the illustrative stub. Even APIs, researchers have found, by the inclusion of an in-practice example of the function described. Spreadsheet programming, then, should not be essentially different: examples of spreadsheet functions in use

V. USER STUDY

Midway through production, we conducted a brief, four-participant user study in order to gauge how well we were capturing the initial design goals. The study was, in general, a flexible and exploratory affair: rather than giving the users any specific tasks or pointed questions, we asked instead that they choose for themselves which trees to explore and report anything that they deemed interesting. Specifically, given trees rooted in the Enron dataset and at least 20 to thirty minutes, the participants situated themselves within the persona of a consultant asked to evaluate a company's spreadsheet practice, which could be arranged around questions like which functions on which the employees most depended, which sheets pre-




Fig. 8. LOOKUP with two (6 occurrences) and three (1648) arguments.

sented the most anomalous or dangerous designs, and so on. Interpretation of “spreadsheet practice” can certainly vary from participant to participant, but since we wanted a range of views and backgrounds despite the persona, we readily accepted this interpretive wiggle room.

We clarified, furthermore, that these notes could be directed either at the quality of the data conveyed by the tool or at the tool’s conveyance itself. Since the tool’s philosophy was based in exploration, we decided that it would not have been appropriate to restrict what people were to look for, letting them find interesting stories of formula construction for themselves. Additionally, either kind of comment – on data or on tool – would direct us back to the central evaluation; a user’s silence, or lack of any interesting findings, would tacitly indicate that perhaps the tool does not convey information well enough.

Still, several limitations cropped up during the study that complicate the evaluation. For one thing, though the source of the data was not explicitly stated at the outset of the session, several users studied the file names (unchanged as they were) to deduce its roots in finance, if not Enron itself. Because of these notorious connotations, then, some were inclined from this moment of discovery to orient their exploration around finance-related functions, potentially limiting their findings or perceived users for the tool – though some also said that, without this essential context, they could not have properly explored a dataset in the first place. Comments like these point out the trade-offs of unguided exploration, too: though it might not restrict them to a single purpose, it might, at the same time, not afford them any mindset at all with which to interpret the data. Another potential problem is that the data which users explored was processed and created before we

had fully completed the tool, meaning some of the counts and patterns in the trees were inaccurate. However, we decided that this was not a pressing concern, being that we were primarily with whether any interesting reports could be made at all, not with whether these reports were 100% accurate assessments of reality.

Nevertheless, these user studies were a formative moment in evaluating the tool’s performance, generating a number of directions or corrections in the process. We present an overview of responses below, with the four participants recorded simply as P1 through P4.

A. Positive Responses

The inclusion of formula examples, grounded in and linking back to the originating spreadsheets, proved to be an especial draw to the tool. After all, none of the participants professed themselves to be experts with Excel, their self-reported proficiencies ranging from “fairly familiar” (P2) to “not extremely familiar” (P4); even if they were experts, it would likely be unreasonable to expect an evenly distributed familiarity with the 100+ function trees available. Nevertheless, the list of actual formulae accessible through a node’s double-click allowed the users a way to find examples of formulae with exactly the configuration they wanted. For example, as P1 explored the VLOOKUP function, even though they knew the official API more verbosely explained its signature, they said they could use the tool to collect more examples than the documentation offers.

Furthermore, it offers a simple way of finding examples where certain functions were used together, though a standard text-search tool might yield the same results for a function alone, combinations can require more complex text queries, whereas here, the information is already captured in nodes. [This paragraph contrasts it with previous approaches, but lacks compelling evidence – a lot of “mights”]

The discovery of errors and bad design, whether by purpose or accident, produced the more jarring events of the study, and as such, most users commented on the tool’s ability to locate these problems in spreadsheet design. This was often supported by the distinctive visual signatures of certain bad practices: multiple arguments produce towers of nodes, deep function nesting stretches the tree horizontally, and so on. As such, the participants remarked on the problems quickly after opening: P3, exploring the example for a 27-argument SUM (with SUM’s initial appearance pictured left), diagnosed the situation with either incompetency or just plain weirdness; P2, meanwhile, decided that the spreadsheet’s author “was having a bad day” when writing an AVERAGE with 18 distinct arguments. Aside from these implicit cues, there are also nodes which refer to formula errors,



The detection of bad design, however, is certainly not limited to a visualization tool, and some users remarked that a tool which printed a list of errors could suffice just as well for that purpose. However, they also remarked that this visual interface better supports the case when searching for new and unknown brands of bad smell.

B. Negative Responses

Furthermore, just because a user can exhaustively explore every variation of how a function was used doesn't mean they will learn exactly what it does. P1, in their exploration, examined functions that they hadn't used before, such as KURT and PV, and described precisely what types of arguments it accepted and how many. However, even after exploring the spreadsheets, they could not accurately describe what the functions produced without consulting the official documentation. Whether through reticence of the tool or obscurity of the functions themselves, the exploratory environment is not enough for independent education. However, unlike the previous problem, this can be addressed, perhaps, by guiding the user directly to such documentation from within the tool without violating any design concerns.

A chord diagram illustrating the relationships between the function IF0 and various other Excel functions. The nodes are arranged in two arcs. The left node is labeled IF0. The right nodes include -NUM-, VLOOKUP(), IF(), -STR-, -REF-, +, AVERAGE(), ROUND(), MATCH(), HLOOKUP(), INDEX(), MIN(), OFFSET(), SUM(), MAX(), &, CHOOSE(), -ERROR-, CONCATENATE(), SUMIF(), %, -BOOL-, DATE(), SQRT(), 0, LINEST(), ABS(), YEAR(), IRR(), LN(), DATEVALUE(), INT(), LOOKUP(), NA(), PRODUCT(), TRIM(), COUNTA(), COUNTIF(), DAY(), MONTH(), NORMDIST(), RAND(), and ROUND(). The thickness of the arcs indicates the frequency or strength of the relationship between IF0 and each function.

considering that the SUM is built from 910 nodes and IF, at least 18000 [double-check these with refined data]. P2 corroborated this by explaining how they gravitated toward specific examples over the plane of nodes but complicates it further by also deeming the tool at its best when it shows either every possibility of an side-by-side (for comparison) or nothing on that branch at all, a balance that is hard to strike with node-dense trees.

The difference is, perhaps, too subtle: on the left, the tree contains instances of AND functions with any number of arguments, up to four; on the right, the tree contains only instances with exactly four arguments. In other words, an instance of "AND(A1, B2, C3, D4)" would be found in both trees, while "AND(A1, B2)" would be found in only the left. Though the users naturally did not explicitly report misinterpreting the icons, it became clear in their out-loud thoughts that they viewed each box as a fundamentally different set of possibilities in formula construction rather than just a specific parameter. Though we clarified the meanings when these misunderstandings became apparent, these events

point to a failure in the tool to properly indicate every function with clarity.

VI. LIMITATIONS

Beyond the negative responses in the user study, a few more inherent limitations beset the tool. These largely resulted from trade-offs in which we found no acceptable way to win the best of both worlds, and necessary compromises were made.

Because the data collection depends on POI's formula parser, it affords no leeway or partial information from a formula; it either processes it perfectly or throws it out. As such, the visualization inhibits insight into anything with syntactical errors or third-party functions that can't be evaluated without additional tools. Note, however, that this does not include standard Excel errors like #REF! and #DIV/0, which the parser handles well. This poses more egregious problems, however, when POI does not support the parsing of certain functions, such as EOMONTH (as obscure as it may be), preventing the healthy growth of those trees.

Many functions have an expected order of arguments, but some, like SUM and MATCH, can be reordered in various ways and maintain their value. However, the representation that this tool uses reinforces for all functions that the order is important, even in these exceptions. Future work could design a new family of tree to accommodate these loose cases, however, by bucketing order-insignificant arguments together; this, of course, would still have to address the profusion of co-presented data.

A larger and related problem is the loss of argument combination. For example, figure VI shows the situation of the INDEX tree, wherein each argument position has a few distinct possibilities. While it is easy to break down, for each position, which argument types are most common (ranges, MATCH, and numbers, respectively), the tree is silent on whether MATCH functions are used together in the second and third arguments, among other concerns. Though these argument combinations could be pieced together by manually inspecting the lists of examples, it is nevertheless inconvenient, though not outside the realm of future improvement.

VII. FUTURE WORK

For these user and case studies, we relied solely on Enron's spreadsheets to find our results. However, Enron's focus on finances and energy represents only one context in which people use spreadsheets. Other corpora, like EUSES and Fuse, draw from other domains and industries, and, as Jansen found in his comparison of spreadsheets Enron and EUSES, different companies rely on different functions. Though the use of a single corpora nevertheless suffices for evaluation, we may yet be able to find more interesting applications and uses for the tool in other types of spreadsheet.

Additionally, Excel is only one of many different spreadsheet products available, chosen for this project because of its prevalence. But other spreadsheets, such as Gnumeric⁹ or




Fig. 10. A difficult question: how often is MATCH used in both the second and third arguments?

Google Sheets¹⁰ represent other approaches to spreadsheet systems. Even within identical domains, there's no guarantee that practices and functions will hold across these varied offerings. This, then, warrants further analysis, in which good visualization might serve well.

VIII. CONCLUSION

This paper presents a tool to support the exploration of function combinations throughout a given spreadsheet dataset. By taking an agnostic approach to design quality, it attempts to convey the dataset's spreadsheet practices as is: the common along with the anomalous, and the well-designed along with the odorous. As seen in the user and case studies, this allows it to take on several potential roles at once, such as an environment for search out bad code smells as well as a repository for examples for learning new techniques.

To be sure, we still have a lot of room for improvement. Though the exploratory philosophy fostered certain goals well, such as the ability to discover new problematic smells without knowing beforehand what they are, it also hindered the project in others. Some users, for example, found the trees of popular functions to be too crowded to explore well, or that their intention to use the tool to examine the good or the bad design exclusively was not fully supported by a tool that claimed to know neither. Many of the problems, however, were not essential to this design conflict and will be overcome in future iterations of design by refining icons and improving maneuverability.

Either way, the visualization of these large datasets represents an important step in improving practice overall. By taking these voluminous oceans of data offered in spreadsheet corpora and rendering them comprehensible through images, we seek to draw attention to valuable information that would otherwise go unnoticed.

⁹<http://www.gnumeric.org/>

¹⁰<https://www.google.com/sheets/about/>

ACKNOWLEDGMENT

This material is based upon work supported in whole or in part with funding from the Laboratory for Analytic Sciences (LAS). Any opinions, findings, conclusions, or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the LAS and/or any agency or entity of the United States Government.

REFERENCES

- [1] R. R. Panko and N. Ordway, "Sarbanes-oxley: What about all the spreadsheets?" *arXiv preprint arXiv:0804.0797*, 2008.
- [2] A. J. Ko, R. Abraham, L. Beckwith, A. Blackwell, M. Burnett, M. Erwig, C. Scaffidi, J. Lawrance, H. Lieberman, B. Myers *et al.*, "The state of the art in end-user software engineering," *ACM Computing Surveys (CSUR)*, vol. 43, no. 3, p. 21, 2011.
- [3] B. A. Nardi and J. R. Miller, *The spreadsheet interface: A basis for end user programming*.
- [4] C. Scaffidi, M. Shaw, and B. Myers, "Estimating the numbers of end users and end user programmers," in *Visual Languages and Human-Centric Computing, 2005 IEEE Symposium on*. IEEE, 2005, pp. 207–214.
- [5] S. G. Powell, K. R. Baker, and B. Lawson, "Impact of errors in operational spreadsheets," *Decision Support Systems*, vol. 47, no. 2, pp. 126–132, 2009.
- [6] M. Fisher and G. Rothermel, "The euses spreadsheet corpus: a shared resource for supporting experimentation with spreadsheet dependability mechanisms," in *ACM SIGSOFT Software Engineering Notes*, vol. 30, no. 4. ACM, 2005, pp. 1–5.
- [7] T. Barik, K. Lubick, J. Smith, J. Slankas, and E. Murphy-Hill, "Fuse: a reproducible, extendable, internet-scale corpus of spreadsheets," in *Proceedings of the 12th Working Conference on Mining Software Repositories*. IEEE Press, 2015, pp. 486–489.
- [8] F. Hermans and E. Murphy-Hill, "Enron's spreadsheets and related emails: A dataset and analysis," in *Proceedings of the 37th International Conference on Software Engineering-Volume 2*. IEEE Press, 2015, pp. 7–16.
- [9] F. Hermans, M. Pinzger, and A. van Deursen, "Detecting code smells in spreadsheet formulas," in *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*. IEEE, 2012, pp. 409–418.
- [10] B. Jansen and F. Hermans, "Code smells in spreadsheet formulas revisited on an industrial dataset," in *Software Maintenance and Evolution (ICSME), 2015 IEEE International Conference on*. IEEE, 2015, pp. 372–380.
- [11] E. Aivaloglou, D. Hoepelman, and F. Hermans, "A grammar for spreadsheet formulas evaluated on two large datasets," in *Source Code Analysis and Manipulation (SCAM), 2015 IEEE 15th International Working Conference on*. IEEE, 2015, pp. 121–130.
- [12] B. Jansen, "Enron versus euses: a comparison of two spreadsheet corpora," *arXiv preprint arXiv:1503.04055*, 2015.
- [13] R. Baeza-Yates, B. Ribeiro-Neto *et al.*, *Modern information retrieval*. ACM press New York, 1999, vol. 463.
- [14] T. Igarashi, J. D. Mackinlay, B.-W. Chang, and P. T. Zellweger, "Fluid visualization of spreadsheet structures," in *Visual Languages, 1998. Proceedings. 1998 IEEE Symposium on*. IEEE, 1998, pp. 118–125.
- [15] M. Clermont, *A scalable approach to spreadsheet visualization*, 2003.
- [16] S. Hipfl, "Using layout information for spreadsheet visualization," *arXiv preprint arXiv:0802.3939*, 2008.
- [17] F. Hermans, M. Pinzger, and A. Van Deursen, "Supporting professional spreadsheet users by generating leveled dataflow diagrams," in *Proceedings of the 33rd International Conference on Software Engineering*. ACM, 2011, pp. 451–460.
- [18] F. Hermans, M. Pinzger, and A. van Deursen, "Breviz: Visualizing spreadsheets using dataflow diagrams," *arXiv preprint arXiv:1111.6895*, 2011.
- [19] M. Fowler, *Refactoring: improving the design of existing code*. Pearson Education India, 2009.
- [20] J. Cunha, J. P. Fernandes, H. Ribeiro, and J. Saraiva, "Towards a catalog of spreadsheet smells," in *Computational Science and Its Applications-ICCSA 2012*. Springer, 2012, pp. 202–216.
- [21] A. Asavametha, "Detecting bad smells in spreadsheets," 2012.
- [22] R. Abreu, J. Cunha, J. P. Fernandes, P. Martins, A. Perez, and J. Saraiva, "Smelling faults in spreadsheets," in *2014 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. IEEE, 2014, pp. 111–120.
- [23] F. Hermans, M. Pinzger, and A. v. Deursen, "Detecting and visualizing inter-worksheet smells in spreadsheets," in *Proceedings of the 34th International Conference on Software Engineering*. IEEE Press, 2012, pp. 441–451.
- [24] S. Badame and D. Dig, "Refactoring meets spreadsheet formulas," in *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*. IEEE, 2012, pp. 399–409.
- [25] G. C. Murphy, M. Kersten, and L. Findlater, "How are java software developers using the elipse ide?" *Software, IEEE*, vol. 23, no. 4, pp. 76–83, 2006.
- [26] M. P. Robillard and R. Deline, "A field study of api learning obstacles," *Empirical Software Engineering*, vol. 16, no. 6, pp. 703–732, 2011.